


### Problem Statement

- The **Face verification** is the process of **extracting facial features** from an image and then **comparing** them to the facial features of another image to **verify** the identity of the person in the image.  
Real-life example: **face ID feature** of cell phones.



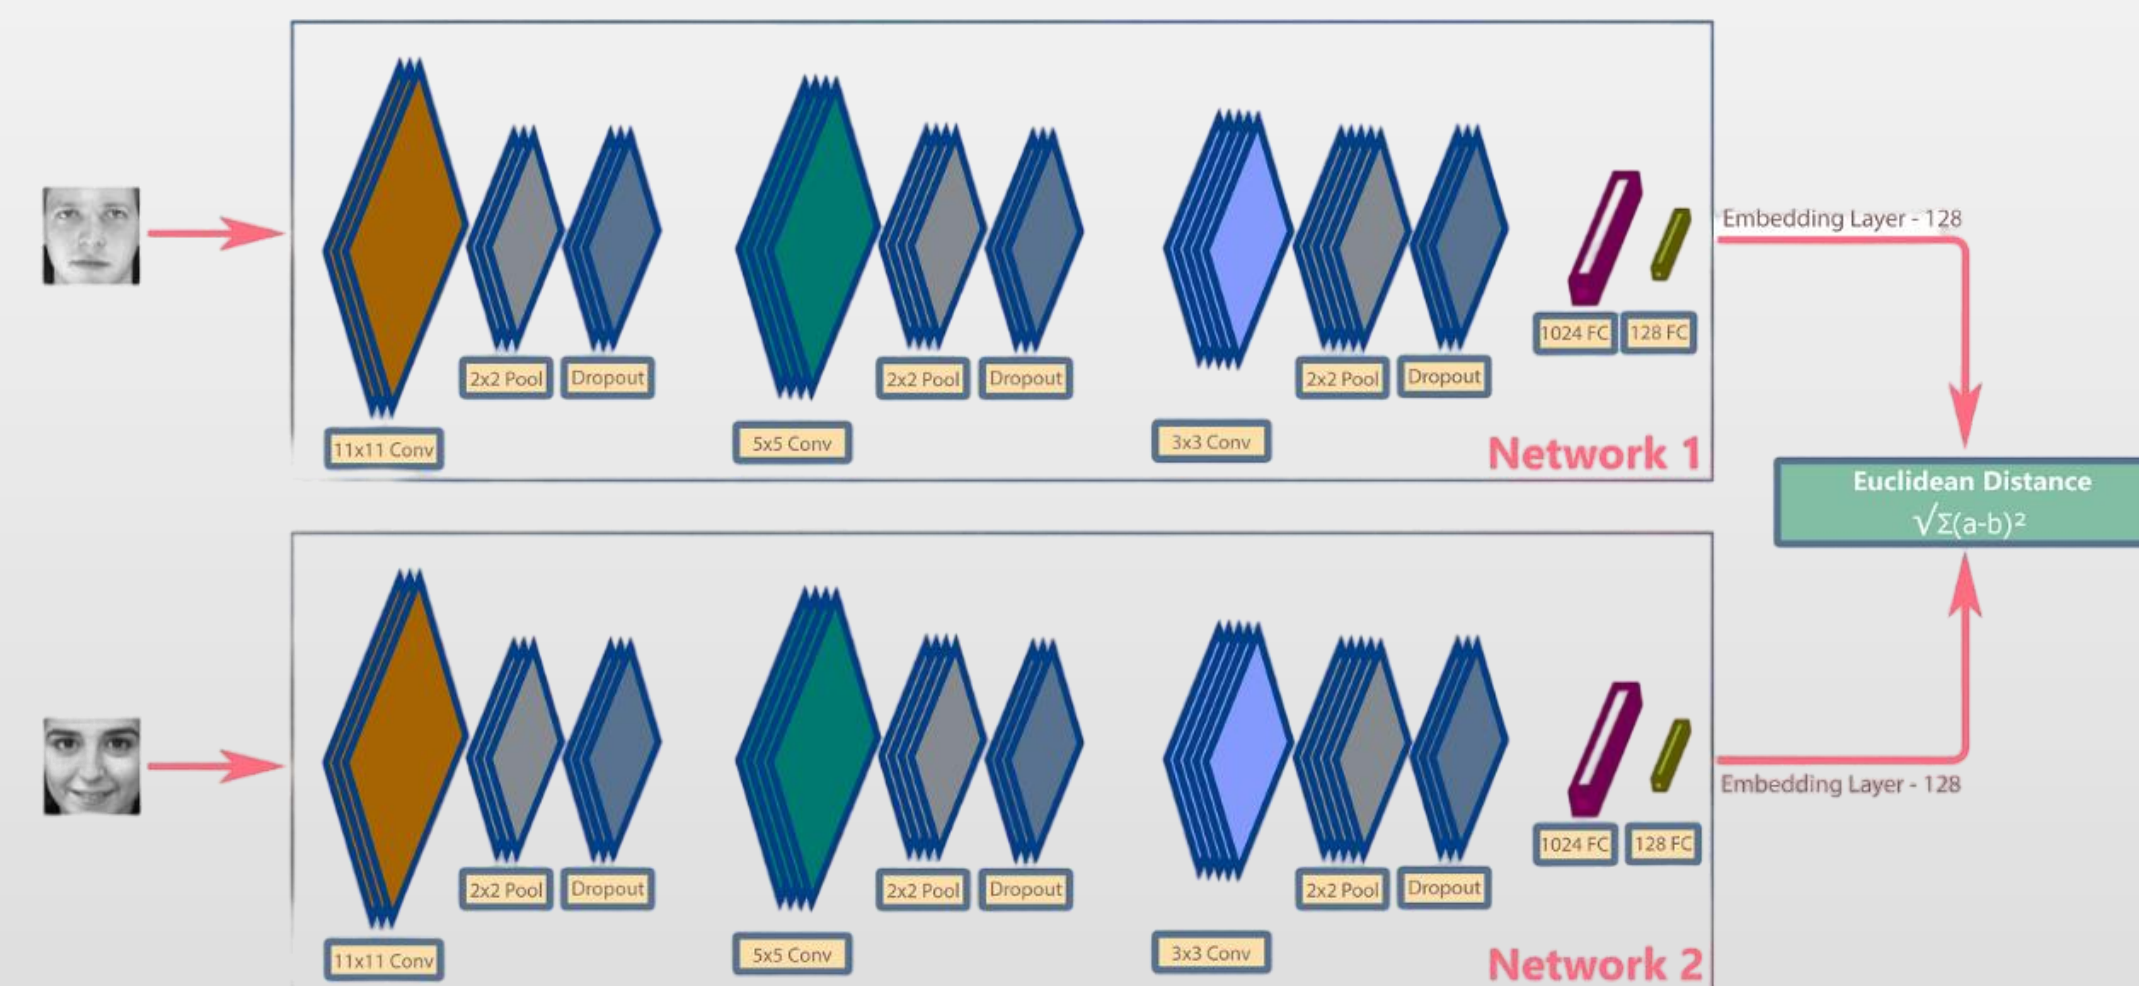
### Dataset Description

- Name: **Extracted Faces, face-recognition-dataset**, derived from the **LFW Dataset**.
- Size: 190 MB
- Description: Each picture is **centered** on a **single face** and encoded in **RGB** and of size **128x128x3**. **1680 celebrity** directories, with **2-50 images** per directory.
- Faces extracted from the original image using **Haar-Cascade Classifier (cv2)**.



### Original chosen model

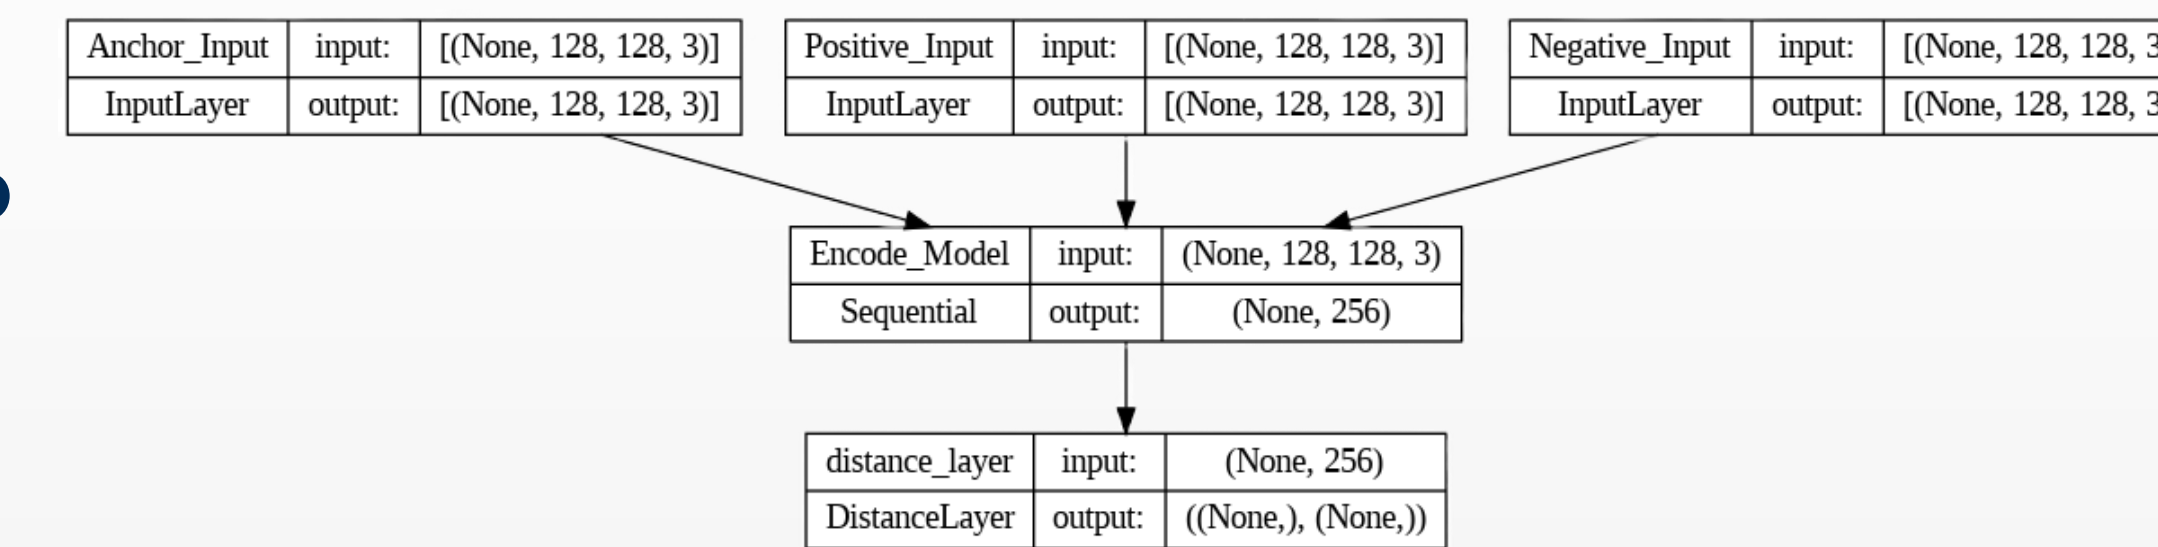

- Siamese Network**: It **does not classify** the images into categories, rather, it finds out the **distance** between any two given images. If the images have the **same label**, then the network should **learn the parameters**, in such a way that it should produce a **smaller distance** between the two images, and if they belong to **different labels**, then the **distance** should be **larger**.



- Encoder**: is responsible for **converting** the images into their **feature vectors**.
- A pre-trained model, **Xception** is used.
- It is connected to **Dense layers** and the last layer uses **L2 Normalization**.

### How it works

- The **Siamese Network** takes 3 input images (**anchor, positive, negative**) and uses the **encoder** to get the **feature vectors**.
- Those features are passed to a **distance layer** which computes the **distance between (anchor, positive) and (anchor, negative) pairs**.

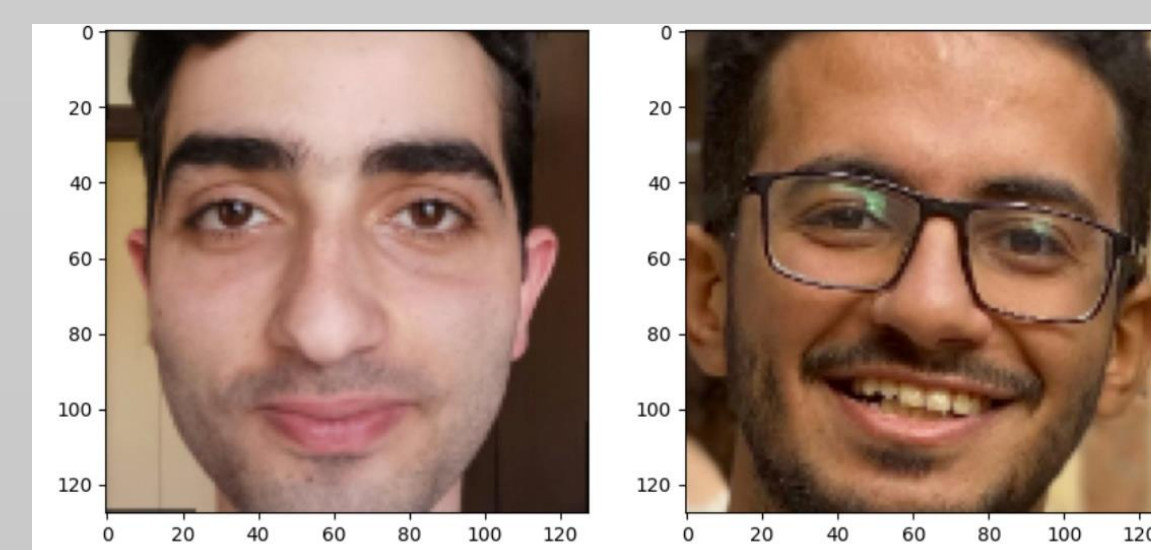


- After training (using the **triplet loss function**), **Encoder** is **extracted** and given two pictures to get their **feature vectors** and calculate the **distance** between them.
- Given a **threshold**, it's determined **whether** the images are the **same person** or a **different person**.

### Updates to the model

- Data is split into correct and fixed sets and triplets.
- Adjusted the number of frozen layers in the Xception model to 35 layers.
- Data augmentation was used to compensate for the small dataset and as a way of regularization.
- Shuffled the training data.
- Added function to do face extraction and alignment.
- Added more evaluation metrics besides accuracy.
- Live face verification.
- Hyperparameter tuning:
  - Tanh instead of Relu
  - Max pooling instead of Average
  - Batch sizes 64, 128, 256
  - Augmentation percentages of 30%, 50%, 60%, 100%

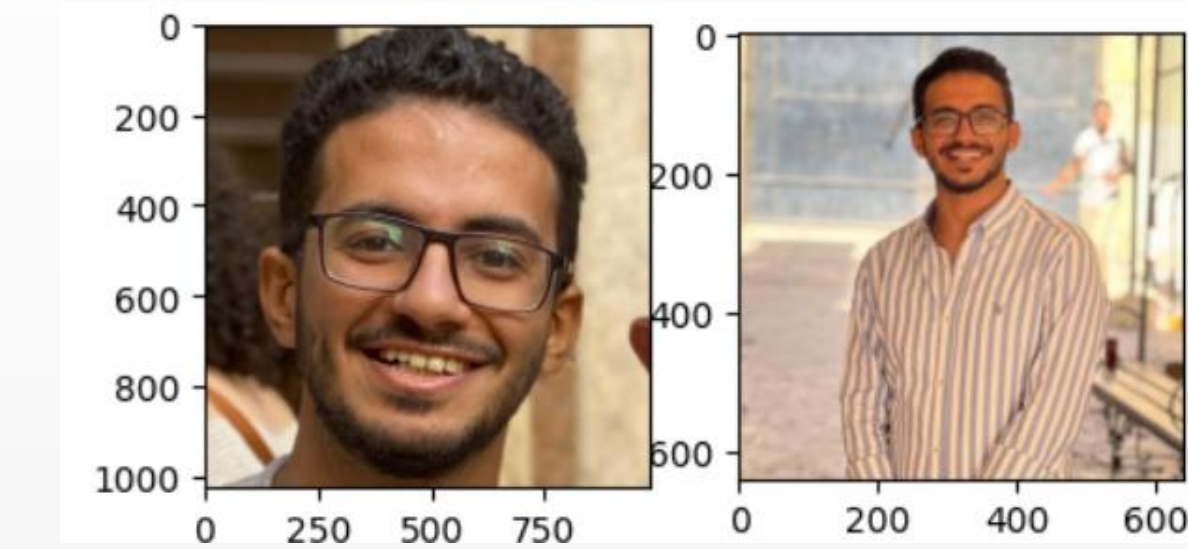
### Input / output example



Distance: [1.7238022]  
prediction is Different

Input

Output




Distance: [0.44238418]  
prediction is Same Person

Input


Output

### Initial Results Vs Final Results

- Initial Accuracy: 84.375%**



- Final Accuracy: 93.164%**
- Recall: 95.7%
- Precision: 91.1%
- F1 score: 93.3%
- AUC: 93.16%
- Categorical Cross-entropy: 2.46



### Conclusion

- Learned to work with limited resources (colab + small dataset) by using the pre-trained model and data augmentation.
- One-Shot Learning.
- Model is biased.
- Results are still far from SOTA or product.

### Acknowledgements

We would like to thank Prof. Moustafa Youssef and Eng. Sherif Mostafa for their continuous support.